FreshWay — Dataset guide

Smart Logistics (EDSLOG.24) · autumn 2026 · public, released in lesson 3

FreshWay's operational data covers **thirteen weeks: ISO week 42 of 2024 through week 2 of 2025** (14 October 2024 – 12 January 2025). That period deliberately contains the busiest week of the year and the quietest one, four calm weeks in between, and one day on which the mechanisation in the central DC failed.

What exists, and who has it

Master data is public. Everyone has it, because no join works without it and because in a real company everyone has master data.

File	What is in it
`stores.csv`	120 stores: format, sales area, coordinates, distance to DC, dock type, unloading window, zero-emission zone yes/no
`product_groups.csv`	20 product groups: temperature zone, shelf life, average price, share of range and revenue
`vehicles.csv`	54 tractor units and **82** delivery vans: type, capacity, powertrain, range, consumption
`suppliers.csv`	~450 suppliers: product groups, EDI connection, agreed lead time and slot window
`sites.csv`	3 DCs, the e-fulfilment centre and 5 hubs: floor area, temperature zones, dock doors, **contracted grid capacity**

Transaction data belongs to a department. You receive your own; the rest exists and you do not have it.

File	Department	What is in it
`inbound_receipts.csv`	**Incoming transport + Storage**	every inbound shipment: supplier, expected and actual arrival, pallets, temperature on arrival, quality check, put-away

File	Department	What is in it
`store_replenishment.csv`	**Outbound transport**	every store delivery: roll cages by temperature zone, time window, planned and actual arrival, vehicle, driving and unloading time, energy, cost
`pick_lines_sample.csv`	**Orderpicking**	pick lines at one regional DC, a documented sample — every twelfth line, and the fraction is in `sample_fraction`: location, quantity, pick time, picker, own or agency staff, errors
`pick_daily_summary.csv`	**Orderpicking + Storage**	throughput per day, per site, per shift, per order type and per temperature zone
`home_delivery_orders.csv`	**E-commerce**	the promise: order time, chosen slot, slot price, lines ordered, substituted and missing, order value, crates
`home_delivery_execution.csv`	**Delivery to consumers**	the execution: actual delivery time, route and stop, vehicle, kilometres, energy, crate temperatures, returns, rating, complaint code
`home_delivery_daily.csv`	**E-commerce + Delivery**	**all six locations**: slots offered, slot capacity, orders delivered and fill rate per day, per site and per two-hour window. Slot capacity is what the OWN fleet was planned to carry, so the fill rate can pass 100
`delivery_routes.csv`	**Delivery to consumers**	every route: vehicle, site, first and last drop, stops, crates, kilometres, energy
`demand_inventory.csv` (plus `_sample` and `_weekly`)	**Inventory and SCM**	per store, per product group, per day: units sold, forecast, stock, days of cover, replenishment, shrink with reason, promotion flag. The weekly file adds `forecast_error_pct`
`site_energy_hourly.csv`	**Storage**	per site, per hour: electricity, peak load, refrigeration, own generation, gas, outside

File	Department	What is in it
		temperature, vehicle charging

The keys — this is where the departments meet

Key	Connects
`store_id`	store replenishment ↔ demand and inventory
`product_group`	inbound ↔ picking ↔ demand and inventory
`order_id`	**the promise ↔ the execution** — and picking, through `order_ref`
`order_ref`	pick lines ↔ store deliveries **or** consumer orders
`vehicle_id`	each flow ↔ `vehicles.csv`. **Two separate fleets**: the trucks in store replenishment and the vans in home delivery share no `vehicle_id`, so this key does not join the two flows to each other
`site_id`	every flow ↔ site energy and contracted capacity
`supplier_id`	inbound receipts ↔ suppliers
`date`	everything

Two files cover one delivery region, not the whole country. `home_delivery_orders.csv` and `home_delivery_execution.csv` describe the orders of **two fulfilment locations** — around 18% of FreshWay's online volume. They are deep, not wide. **National figures come from `home_delivery_daily.csv`**, which covers all six locations. Do not present a number from the order files as a company figure.

Do not trust a derived column without checking its definition. `on_time` in the execution file counts a delivery as on time if it arrived before the window closed — including deliveries that arrived *before the window opened*. "Inside the promised slot" is a different number, and you have both timestamps to work it out yourself. There are more columns like this; the data dictionary tells you what was calculated and how.

One warning before you join orders to execution. The two files do not line up one-to-one: about **18% of online orders are click & collect** and are never delivered by van, so they appear in `home_delivery_orders.csv` and not in `home_delivery_execution.csv`. An inner join silently drops

them. That is not a fault in the data — it is a fact about the company, and if your order count changes when you join, this is why.

And one about complaints. Complaints and ratings are registered at the door, so they sit with Delivery to consumers. Not all of them are caused there: a missing item was decided before the van left. The complaint code tells you the cause; proving how the causes divide needs the order file as well as the execution file, and those sit with two departments.

Two things follow from this table, and they are the point of the whole set.

Not every question can be answered from your own data. Some can: your own volumes, your own costs, your own delays. Others cannot, by construction. "Did we keep our delivery promise?" needs the promise *and* the execution, and those sit with two different departments. So does "what did that late supplier actually cost us" — that one runs across three.

Nothing is hidden from you that a colleague cannot give you. If you need a set, ask the department that owns it. That is what the Data Council is for, and an Interface Agreement is how you record what you agreed. Do not file an Information Request with the Board for something a colleague already has: the Board will refuse and send you to them.

How to work with it

- **Read week 52 by the day, not by the week.** Two days in that week carry more than half of it and two days are almost empty. A weekly average will tell you the opposite of what happened.
- **Large files come in three forms**: the full set, an extract of around 60,000 rows for Excel, and a daily aggregate. Use the aggregate for the overview and the full set for the analysis.
- **A sampled file says so in a column.** `pick_lines_sample.csv` carries `sample_fraction`; divide by it before you draw a conclusion about volumes.
- **Empty is not zero.** A missing temperature reading is not 0 °C, and a missing rating is not a bad one.
- **Money.** Cost fields are excluding VAT; `order_value_eur` and `revenue_eur` are including VAT, because that is what the customer pays.
- Every set has a `data_dictionary.csv` next to it. Read it before you start guessing what a column means.

> **GenAI transparency.** This guide was written with Claude (Anthropic) by the course team as part of the FreshWay case design. The datasets it describes are generated from a source-checked figure sheet; no figures were invented by the model. Human-AI Agency Label: **Co-Creator**. Label source: Boetje, J.,

> & Baake, G. (2026). *Nine prototypical human-AI agency patterns* [Figure]. figshare. https://doi.org/10.6084/m9.figshare.31706884 (CC BY 4.0).

Levering · gerenderd uit `Dataset guide.md` · 2026-09-02